A cost and information-based account of epistemic *must*

Judith Degen, Justine Kao, Gregory Scontras, Noah D. Goodman (Stanford University)
jdegen@stanford.edu

We show how a model of rational inference in communication delivers the puzzlingly weak interpretation of the necessity modal *must*. How could *must q* (1b) not entail that *q* (1a)?

(1a) It's raining.

(1b) It must be raining.

Since Karttunen (1972), linguists have debated the meaning of *must*, arguing about its semantic strength. Rather than engineering weakness into the meaning of the word *must*, our account derives its weakness as an M-implicature (Levinson, 2000): *must q* is marked (i.e., costly) relative to the bare form (1a); the bare form is sufficiently strong already to convey that it is raining (*q*), so listeners take the marked form to convey the marked meaning that the speaker arrived at the conclusion *q* via an evidentially less certain route than if they had chosen the shorter bare form.

To test this account, we determine empirically whether a) a speaker's choice between (1a) and (1b) is affected by the strength of evidence for *q*; b) listeners' interpretations of (1a) and (1b) differ with respect to the strength of their resulting belief in *q*; and c) these beliefs are determined partly by the evidence they attribute to the speaker's choice between (1a) and (1b).

In **Exp.~1 (n=40)**, we collected estimates of evidence strength. Participants on Amazon's Mechanical Turk rated the probability of *q* (e.g., of rain) given a piece of evidence (e.g., *You hear the sound of water dripping on the roof*) on a sliding scale with endpoints labeled ``impossible" and ``certain". These estimates were used for analysis in Exps. 2 and 3.

Exp.~2 (n=40) tested how likely speakers are to use the marked *must p* utterance as evidence strength decreases. On each trial, participants were presented with a piece of evidence (e.g., *You see a person come in from outside with wet hair and wet clothes*) and were asked to choose one of four utterances – bare (1a), *must p* (1b), *probably p*, *might p* – to describe the situation to a friend. Participants were more likely to choose the more marked *must* form over the bare form as the strength of evidence decreased (β=5.4, SE=2.4, p<.05), even when controlling for evidence type (e.g., perceptual, reportative, inferential).

Exp.~3 (n=120) tested whether listeners' estimates of a) the probability of *q* and b) the strength of speakers' evidence for *q* differ depending on the observed utterance; i.e. whether listeners take into account their knowledge of speakers' likely utterances in different evidential states as they interpret the bare and *must* forms. On each trial, participants saw an utterance (e.g. *It's raining*), and were asked a) to rate the probability of *q* on a sliding scale with endpoints labeled ``impossible" and ``certain"; and b) to select one out of five pieces of evidence that the speaker must have had about *q*. Participants believed *q* was less likely after observing the *must* utterance (μ=.65, sd=.21) than after observing the bare utterance (μ=.86, sd=.15, β =-.21, SE=.02, t=-10.1, p<.0001). In addition, average strength of evidence was lower after *must* (μ=.78, sd=.12) than after the bare utterance (μ=.87, sd=.1, β =-.08, SE=.01, t=-6.8, p<.0001).

Following Lassiter and Goodman (2013) we present an extension of the Bayesian Rational Speech Act framework (Frank and Goodman, 2012) using lexical uncertainty to derive the implicature. In this model, the semantics of the bare utterance and *must q* are relatively unconstrained. We define the semantics of the utterances such that $p(q|bare) > \theta_b$ and $p(q|must) > \theta_m$, where the pragmatic listener is uncertain about θ_b and θ_m and infers the values through pragmatic reasoning. When the cost of uttering *must q* is greater than the bare form, the pragmatic listener infers that $p(q)$ is smaller than when the utterance is the less costly bare *q*. Given the weakened certainty of *q*, the listener may then infer that the speaker has weak or imperfect evidence of *q*. Our empirical results and computational model support this account and provide a new perspective on the meaning of *must*: its weakened meaning derives straightforwardly from an M-implicature. We discuss this model with respect to considerations of efficiency in the tradeoff of production and comprehension costs.